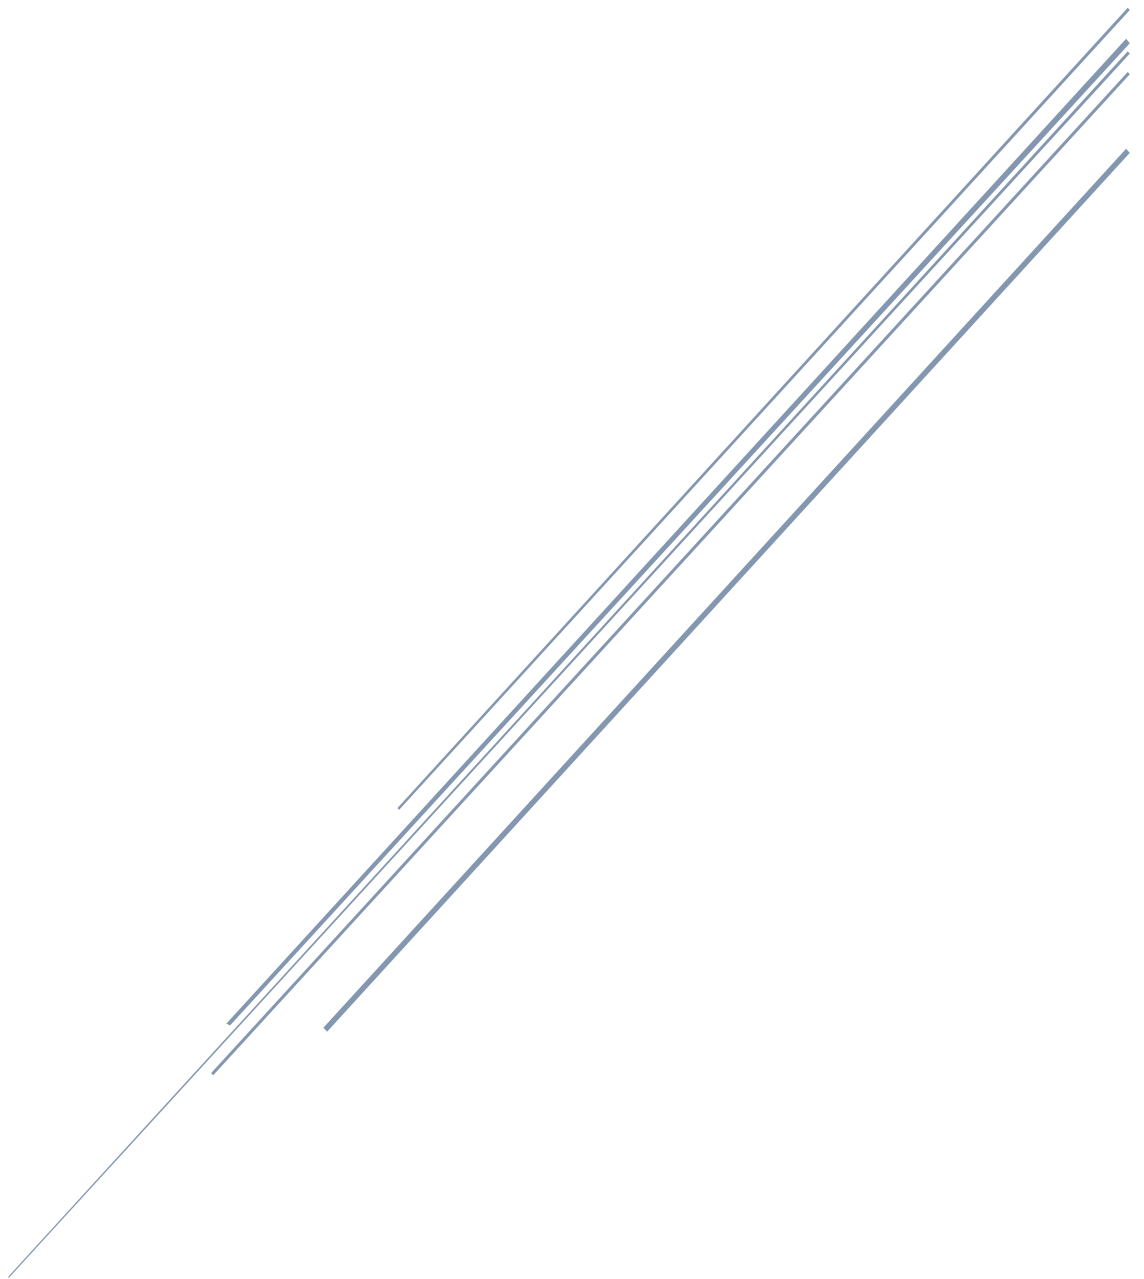


DAT151 – OBLIG5

Optimalisering og implementasjon av større database.

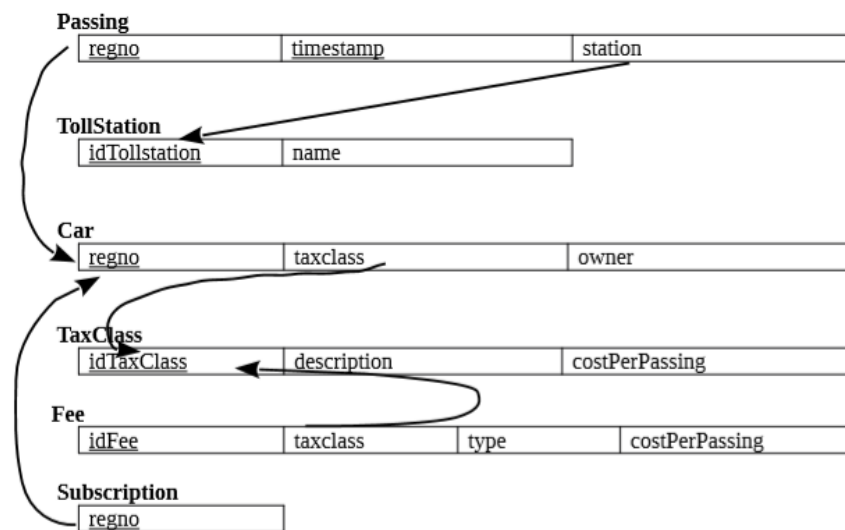


Adrian R.J. Mortensen

Innhold

Oppgave 1: Normal form	2
1. Er skjemaet 1NF?.....	2
2. Er det 2NF?	2
3. er den 3nf	2
Oppgave 2: Implementasjon av fysisk skjema og test miljø.....	3
Oppgave 3: Optimalisering av database	6
a) Query 1	6
b) Query 2	11
c) Query 3	14
d) Query 4 (same as 2).....	16
e) Query 5 (same as 3).....	16
f) Query 5	18

Oppgave 1: Normal form



1. Er skjemaet 1NF?

Et skjema er 1NF hvis og bare hvis alle underliggende domener har atomiske verdier.

All data er atomære om vi antar «timestamp» er det.

2. Er det 2NF?

Et skjema er 2NF hvis og bare hvis det er 1NF og alle ikke nøkkel attributter er knyttet til en kandidat nøkkel eller en annen ikke nøkkel attributt.

Her kan man diskutere at Passing ikke følger denne regelen. Ettersom tollstasjon ikke er avhengig av hverken «regno» eller «timestamp».

3. er den 3nf

Om man mener station er knyttet til timestamp og regno er den det. Ellers er den ikke det siden den ikke fyller 2nf.

Oppgave 2: Implementasjon av fysisk skjema og test miljø.

Jeg bruker skjemaet fra forrige oppgave uten endringer.

SQL filer kan finnes på <https://github.com/H571531/DAT151/tree/master/Oblig5>

> Source /home/admo/git/DAT151/Oblig5/SQL/CreateTable.sql

```
CREATE TABLE IF NOT EXISTS ImportTable (
    regno VARCHAR(45) NULL,
    tid DATETIME NULL,
    idTollstation VARCHAR(45) NULL,
    tollname VARCHAR(45) NULL,
    OwnerName VARCHAR(45) NULL,
    taxId varchar(3),
    taxDesc VARCHAR(45) NULL,
    Subscription VARCHAR(10),
    TollSFee INT(11) NULL,
    SubFee INT(11) NULL
)
ENGINE = MyISAM;

CREATE TABLE IF NOT EXISTS Tollstation (
    idTollstation SMALLINT UNSIGNED NOT NULL,
    name VARCHAR(85) NULL,
    PRIMARY KEY ( idTollstation ))
ENGINE = InnoDB;

CREATE TABLE IF NOT EXISTS TaxClass (
    idTaxClass SMALLINT UNSIGNED NOT NULL,
    description TEXT NULL,
    PRIMARY KEY ( idTaxClass )
)
ENGINE = InnoDB;

CREATE TABLE IF NOT EXISTS Fee (
    idFee SMALLINT NOT NULL AUTO_INCREMENT,
    taxclass SMALLINT UNSIGNED NOT NULL,
    type ENUM('regular', 'withsubscription') NOT NULL,
    costPerPassing DECIMAL(5,2) NOT NULL,
    PRIMARY KEY ( idFee ),
    CONSTRAINT TeacherFK FOREIGN KEY (taxclass) REFERENCES
TaxClass(idTaxClass)
)
ENGINE = InnoDB;

CREATE TABLE IF NOT EXISTS Car (
    regno CHAR(7) NOT NULL,
    owner VARCHAR(85) NULL,
    taxclass SMALLINT UNSIGNED NOT NULL,
    PRIMARY KEY ( regno ),
    CONSTRAINT fk_Car_TaxClass1
    FOREIGN KEY ( taxclass )
    REFERENCES TaxClass (idTaxClass)
    ON DELETE NO ACTION
    ON UPDATE NO ACTION)
ENGINE = InnoDB;
```

```

CREATE TABLE IF NOT EXISTS Passing (
    regno CHAR(7) NOT NULL,
    timestamp TIMESTAMP NOT NULL,
    tollstation SMALLINT UNSIGNED NOT NULL,
    PRIMARY KEY ( regno , timestamp ),
    CONSTRAINT fk_Passing_1
        FOREIGN KEY ( tollstation )
        REFERENCES Tollstation ( idTollstation )
        ON DELETE NO ACTION
        ON UPDATE NO ACTION,
    CONSTRAINT fk_Passing_Subscription1
        FOREIGN KEY ( regno )
        REFERENCES Car ( regno )
        ON DELETE NO ACTION
        ON UPDATE NO ACTION)
ENGINE = InnoDB;

```

```

CREATE TABLE IF NOT EXISTS Subscription (
    regno CHAR(7) NOT NULL,
    PRIMARY KEY ( regno ),
    CONSTRAINT fk_Subscription_Car1
        FOREIGN KEY ( regno )
        REFERENCES Car ( regno )
        ON DELETE NO ACTION
        ON UPDATE NO ACTION)
ENGINE = InnoDB;

```

Get data

```
#!/bin/bash
```

```
wget https://eple.hib.no/fag/dat151/v2020/carpassingdb.txt
```

(reason for making a script is due to Github limitations (Can't push files over 100MB))

```

> LOAD DATA LOCAL INFILE '/home/admo/git/DAT151/Oblig5/SQL/carpassingdb.txt' INTO TABLE
ImportTable FIELDS TERMINATED BY ',';

```

Move data to tables

```
-- Tollstation
INSERT INTO Tollstation
SELECT DISTINCT idTollstation,tollname
FROM ImportTable;

-- TaxClass
INSERT INTO TaxClass
SELECT DISTINCT taxId,taxDesc
FROM ImportTable;

-- Car
INSERT INTO Car
SELECT DISTINCT regno,OwnerName,taxId
FROM ImportTable;

-- Subscription
INSERT INTO Subscription
SELECT DISTINCT regno
FROM ImportTable
WHERE Subscription='yes';

-- Passing
INSERT INTO Passing
SELECT DISTINCT regno,tid,idTollstation
FROM ImportTable;

-- Fee
INSERT INTO Fee (taxclass,type,costPerPassing)
SELECT DISTINCT taxId,'regular',TollsFee
FROM ImportTable
WHERE Subscription='no';

INSERT INTO Fee (taxclass,type,costPerPassing)
SELECT DISTINCT taxId,'withsubscription',SubFee
FROM ImportTable
WHERE Subscription='yes';
```

Oppgave 3: Optimalisering av database

a) Query 1

Finner navnet og tiden til noen som har kjørt igjennom en bomstasjon en spesifikk dag.

```
MariaDB [Oblig5]> SELECT SQL_NO_CACHE C.owner, P.timestamp
-> FROM Car C JOIN Passing P USING(regno)
-> WHERE YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=3
-> AND DAYOFWEEK(P.timestamp)=1;
```

owner	timestamp
Urfan SandÅ,y	2018-03-04 04:30:00
Urfan SandÅ,y	2018-03-04 05:42:00
Birte Fossum	2018-03-04 02:18:00
Birte Fossum	2018-03-04 23:31:00
Stanley Ingvaldsen	2018-03-04 02:18:00
Stanley Ingvaldsen	2018-03-04 23:31:00
Amar Wiig	2018-03-04 05:53:00
Minda Larssen	2018-03-04 14:24:00
Storm Nordstrand	2018-03-04 03:09:00
Marcus Hafstad	2018-03-04 14:24:00
Mathilde Lillevik	2018-03-04 04:30:00
Mathilde Lillevik	2018-03-04 05:42:00
Zilan Solbakken	2018-03-04 05:53:00
Kurt Aslaksen	2018-03-04 19:02:00
Dilara Skar	2018-03-04 03:09:00
Annette Sara	2018-03-04 19:02:00

16 rows in set (3.085 sec)

Profiling before optimization:

Ran the query 10 times before profiling start.

MariaDB [Oblig5]> SHOW PROFILES;

Query_ID	Duration	Query
1	3.08046248	SELECT SQL_NO_CACHE C.owner, P.timestamp FROM Car C JOIN Passing P USING(regno) WHERE YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=3 AND DAYOFWEEK(P.timestamp)=1
2	3.14813736	SELECT SQL_NO_CACHE C.owner, P.timestamp FROM Car C JOIN Passing P USING(regno) WHERE YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=3 AND DAYOFWEEK(P.timestamp)=1
3	3.15549203	SELECT SQL_NO_CACHE C.owner, P.timestamp FROM Car C JOIN Passing P USING(regno) WHERE YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=3 AND DAYOFWEEK(P.timestamp)=1
4	3.16799906	SELECT SQL_NO_CACHE C.owner, P.timestamp FROM Car C JOIN Passing P USING(regno) WHERE YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=3 AND DAYOFWEEK(P.timestamp)=1
5	3.16238553	SELECT SQL_NO_CACHE C.owner, P.timestamp FROM Car C JOIN Passing P USING(regno) WHERE YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=3 AND DAYOFWEEK(P.timestamp)=1

Picking Query 3 for further info:

MariaDB [Oblig5]> SHOW PROFILE for QUERY 3;

Status	Duration
Starting	0.000083
Checking permissions	0.000004
Opening tables	0.000032
After opening tables	0.000006
System lock	0.000002
Table lock	0.000005
Init	0.000029
Optimizing	0.000018
Statistics	0.000021
Preparing	0.000018
Executing	0.000002
Sending data	3.155202
End of update loop	0.000017
Query end	0.000003
Commit	0.000004
Closing tables	0.000003
Unlocking tables	0.000001
Closing tables	0.000010
Starting cleanup	0.000002
Freeing items	0.000008
Updating status	0.000021
Reset for next command	0.000003

Most of the time in this query is spent between Executing and Sending data. (The Duration label is abit misleading...) Each time means the time elapsed between the previous event and the new event.

EXPLAIN before Optimization

MariaDB [Oblig5]> EXPLAIN SELECT SQL_NO_CACHE C.owner, P.timestamp FROM Car C JOIN Passing P USING(regno) WHERE YEAR (P.timestamp)=2018 AND MONTH (P.timestamp)=3 AND DAYOFWEEK(P.timestamp)=1;

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	C	ALL	PRIMARY	NULL	NULL	NULL	203998	
1	SIMPLE	P	ref	PRIMARY	PRIMARY	7	Oblig5.C.regno	1	Using where

2 rows in set (0.000 sec)

Optimization of query and table.

By breaking normalization I might get a better result.

The easiest change I can see right now is remove the need to join the car table with Passing.

The only thing we need from the Car table is the name. Therefore I add name to the Passing.

```
CREATE TABLE IF NOT EXISTS Passing2 (
  regno CHAR(7) NOT NULL,
  timestamp TIMESTAMP NOT NULL,
  tollstation SMALLINT UNSIGNED NOT NULL,
  owner VARCHAR(85),
  PRIMARY KEY ( regno , timestamp ),
  CONSTRAINT fk_Passing_2
    FOREIGN KEY ( tollstation )
    REFERENCES Tollstation ( idTollstation )
    ON DELETE NO ACTION
    ON UPDATE NO ACTION,
  CONSTRAINT fk_Passing_Subscription2
    FOREIGN KEY ( regno )
    REFERENCES Car ( regno )
    ON DELETE NO ACTION
    ON UPDATE NO ACTION)
ENGINE = InnoDB;

-- Passing2
INSERT INTO Passing2
SELECT DISTINCT regno,tid,idTollstation,OwnerName
FROM ImportTable;

-- Query
SELECT SQL_NO_CACHE owner,timestamp
FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3
AND DAYOFWEEK(timestamp)=1;
```

```
+-----+-----+
| owner                | timestamp                |
+-----+-----+
| Urfan SandÃ,y        | 2018-03-04 04:30:00 |
| Urfan SandÃ,y        | 2018-03-04 05:42:00 |
| Birte Fossum         | 2018-03-04 02:18:00 |
| Birte Fossum         | 2018-03-04 23:31:00 |
| Stanley Ingvaldsen   | 2018-03-04 02:18:00 |
| Stanley Ingvaldsen   | 2018-03-04 23:31:00 |
| Amar Wiig            | 2018-03-04 05:53:00 |
| Minda Larssen        | 2018-03-04 14:24:00 |
| Storm Nordstrand     | 2018-03-04 03:09:00 |
| Marcus Hafstad       | 2018-03-04 14:24:00 |
| Mathilde Lillevik    | 2018-03-04 04:30:00 |
| Mathilde Lillevik    | 2018-03-04 05:42:00 |
| Zilan Solbakken       | 2018-03-04 05:53:00 |
| Kurt Aslaksen        | 2018-03-04 19:02:00 |
| Dilara Skar          | 2018-03-04 03:09:00 |
| Annette Sara         | 2018-03-04 19:02:00 |
+-----+-----+
16 rows in set (2.555 sec)
```

Can already see some time saved.

Profiling after optimization

Query has been ran 10 times before profiling was turned on.

```
MariaDB [Oblig5]> show profiles;
```

Query_ID	Duration	Query
1	2.46529822	SELECT SQL_NO_CACHE owner,timestamp FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3 AND DAYOFWEEK(timestamp)=1
2	2.51719002	SELECT SQL_NO_CACHE owner,timestamp FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3 AND DAYOFWEEK(timestamp)=1
3	2.52976462	SELECT SQL_NO_CACHE owner,timestamp FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3 AND DAYOFWEEK(timestamp)=1
4	2.51797654	SELECT SQL_NO_CACHE owner,timestamp FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3 AND DAYOFWEEK(timestamp)=1
5	2.47903186	SELECT SQL_NO_CACHE owner,timestamp FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3 AND DAYOFWEEK(timestamp)=1

5 rows in set (0.000 sec)

Showing for query 3

```
MariaDB [Oblig5]> show profile for query 3;
```

Status	Duration
Starting	0.000077
Checking permissions	0.000005
Opening tables	0.000019
After opening tables	0.000004
System lock	0.000003
Table lock	0.000006
Init	0.000025
Optimizing	0.000014
Statistics	0.000014
Preparing	0.000019
Executing	0.000002
Sending data	2.529506
End of update loop	0.000018
Query end	0.000002
Commit	0.000005
Closing tables	0.000003
Unlocking tables	0.000001
Closing tables	0.000009
Starting cleanup	0.000002
Freeing items	0.000006
Updating status	0.000021
Reset for next command	0.000003

22 rows in set (0.000 sec)

Explain optimized query

```
MariaDB [Oblig5]> EXPLAIN SELECT SQL_NO_CACHE owner,timestamp  
-> FROM Passing2 WHERE YEAR(timestamp)=2018 AND MONTH(timestamp)=3  
-> AND DAYOFWEEK(timestamp)=1;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	Passing2	ALL	NULL	NULL	NULL	NULL	5035359	Using where

1 row in set (0.001 sec)

Indexes

Now attempting to add a index on the timestamp collumn to perhaps make the query abit faster.

```
CREATE INDEX Timestamp ON Passing2(timestamp);
```

This adds very little to making the query any faster.

This also counts for Passing (the original table) the optimizer does not even want to use the index.

My thoughts are that this is because there is no need for it.

Note:

It is worth noting that on pre cache queries The un optimized query used 7s and the optimized query used 4s.

b) Query 2

Finner de som har abonnement med total kostnad over 4000

```
CREATE INDEX Timestamp ON Passing(timestamp);
-- b
MariaDB [Oblig5]> SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
-> FROM Car C JOIN Passing P USING(regno)
-> JOIN TaxClass T ON C.taxclass=T.idTaxClass
-> JOIN Fee F ON F.taxclass=T.idTaxClass
-> JOIN Subscription S USING(regno)
-> WHERE F.type='withsubscription'
-> GROUP BY C.owner HAVING totalfee > 4000;
```

```
+-----+-----+
| carowner | totalfee |
+-----+-----+
| Ansgar Oftedal | 4210.00 |
| Rina Kvalheim | 4605.00 |
+-----+-----+
2 rows in set (2.475 sec)
```

```
MariaDB [Oblig5]> EXPLAIN SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
-> FROM Car C JOIN Passing P USING(regno)
-> JOIN TaxClass T ON C.taxclass=T.idTaxClass
-> JOIN Fee F ON F.taxclass=T.idTaxClass
-> JOIN Subscription S USING(regno)
-> WHERE F.type='withsubscription'
-> GROUP BY C.owner HAVING totalfee > 4000;
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | T | index | PRIMARY | PRIMARY | 2 | NULL | 10 | Using index; Using temporary; Using filesort |
| 1 | SIMPLE | F | ref | TeacherFK | TeacherFK | 2 | Oblig5.T.idTaxClass | 1 | Using where |
| 1 | SIMPLE | C | ref | PRIMARY, fk_Car_TaxClass1 | fk_Car_TaxClass1 | 2 | Oblig5.T.idTaxClass | 1 | |
| 1 | SIMPLE | S | eq_ref | PRIMARY | PRIMARY | 7 | Oblig5.C.regno | 1 | Using index |
| 1 | SIMPLE | P | ref | PRIMARY, reg | PRIMARY | 7 | Oblig5.C.regno | 13 | |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
5 rows in set (0.001 sec)
```

The optimizer already uses indexes and I cant see a better way to do the indexes. However I want to attempt denormalization to see if it can get better by denormalization.

I see that Taxclass is joined with the cars taxclass. Only to get fee. Now if we get the fee into passing we wont have to do those two joins. And if we also add owner to the table we might just have to find the cars with subscription. Adding who has subscription aswell might be faster. Butt his will increase the column count with 3 so the negative effects of this could be negative to the outcome.

Trying this either way.

Optimizing

Sql: Table

```
CREATE TABLE IF NOT EXISTS Passing3 (
  regno CHAR(7) NOT NULL,
  timestamp TIMESTAMP NOT NULL,
  tollstation SMALLINT UNSIGNED NOT NULL,
  owner VARCHAR(85),
  type ENUM('regular', 'withsubscription') NOT NULL,
  costPerPassing DECIMAL(5,2) NOT NULL,
  PRIMARY KEY ( regno , timestamp ),
  CONSTRAINT fk_Passing_3
    FOREIGN KEY ( tollstation )
    REFERENCES Tollstation ( idTollstation )
    ON DELETE NO ACTION
    ON UPDATE NO ACTION,
  CONSTRAINT fk_Passing_Subscription3
    FOREIGN KEY ( regno )
    REFERENCES Car ( regno )
    ON DELETE NO ACTION
    ON UPDATE NO ACTION)
ENGINE = InnoDB;
```

Adding data:

```
-- Passing3
INSERT INTO Passing3
SELECT DISTINCT regno,tid,idTollstation,OwnerName,'withsubscription',SubFee
FROM ImportTable
WHERE Subscription='yes';

INSERT INTO Passing3
SELECT DISTINCT regno,tid,idTollstation,OwnerName,'regular',TollSFee
FROM ImportTable
WHERE Subscription='no';
```

Query:

```
SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee
from Passing3
WHERE type='withsubscription'
GROUP BY owner HAVING totalfee > 4000;
```

We can see the Query being alot shorter with no joins and maybe arguably easier to read.

However did it increase performance?

Profiling

Pre denormalization

MariaDB [Oblig5]> show profiles;

Query_ID	Duration	Query
1	3.51759643	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass JOIN Subscription S USING(regno) WHERE F.type='withsubscription'
2	3.53729194	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass JOIN Subscription S USING(regno) WHERE F.type='withsubscription'
3	3.51002205	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass JOIN Subscription S USING(regno) WHERE F.type='withsubscription'
4	3.47434126	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass JOIN Subscription S USING(regno) WHERE F.type='withsubscription'
5	3.39511616	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass JOIN Subscription S USING(regno) WHERE F.type='withsubscription'

5 rows in set (0.000 sec)

Post denormalization

MariaDB [Oblig5]> show profiles;

Query_ID	Duration	Query
1	2.65984540	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='withsubscription'
2	2.66344669	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='withsubscription'
3	2.66824841	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='withsubscription'
4	2.67204567	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='withsubscription'
5	2.66805130	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='withsubscription'

5 rows in set (0.000 sec)

Yes! The performance seem to have been improved. Only with about a second. Both profiles started after the query was ran 10 times.

Having a look at the explain for fun to see how this new query is being optimized by the optimizer.

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	Passing3	ALL	NULL	NULL	NULL	NULL	4812140	Using where; Using temporary; Using filesort

1 row in set (0.001 sec)

From the explain we can see there is no joins, and no other tables to take into consideration. If there was; the lack of keys could mean we would need some sort of indexes. But we only have a select on a single table. And since we're using where clauses I believe this is sufficiently improved from the original query.

c) Query 3

This query does the same as the previous one. This time using joins.

```
MariaDB [Oblig5]> SELECT SQL_NO_CACHE C.owner AS carowner,  
-> Sum(F.costPerPassing) AS totalfee  
-> FROM Car C JOIN Passing P USING(regno)  
-> JOIN TaxClass T ON C.taxclass=T.idTaxClass  
-> JOIN Fee F ON F.taxclass=T.idTaxClass  
-> WHERE F.type='withsubscription'  
-> AND C.regno IN (SELECT regno FROM Subscription)  
-> GROUP BY C.owner HAVING totalfee > 4000;
```

```
+-----+-----+  
| carowner          | totalfee |  
+-----+-----+  
| Ansgar Oftedal   | 4210.00 |  
| Rina Kvalheim    | 4605.00 |  
+-----+-----+  
2 rows in set (3.330 sec)
```

Explain

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	PRIMARY	T	index	PRIMARY	PRIMARY	2	NULL	10	Using index; Using temporary; Using filesort
2	PRIMARY	F	ref	TeacherFK	TeacherFK	2	Oblig5.T.idTaxClass	1	Using where
3	PRIMARY	C	ref	PRIMARY_fk_Car_TaxClass1	fk_Car_TaxClass1	2	Oblig5.T.idTaxClass	1	Using where
4	PRIMARY	Subscription	eq_ref	PRIMARY	PRIMARY	7	Oblig5.C.regno	1	Using index
5	PRIMARY	P	ref	PRIMARY_reg	PRIMARY	7	Oblig5.C.regno	13	

5 rows in set (0.001 sec)

Here we can see we have available keys on all joins. This is good. The Optimizer will be using indexes some places. The Key length is also fairly small meaning it only needs a few bytes of the key. Therefore I will not be addressing any optimization by indexes.

Profiling

MariaDB [Oblig5]> show profiles;

Query_ID	Duration	Query
1	3.00203704	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.owner HA
2	3.05971641	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.owner HA
3	3.00294968	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.owner HA
4	3.11308706	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.owner HA
5	3.08940385	SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T ON C.taxclass=T.idTaxClass JOIN Fee F ON F.taxclass=T.idTaxClass WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.owner HA

5 rows in set (0.000 sec)

In profiling we can see that this Query is performs better than the previous one that used joins instead of subqueries. But still not as good as the denormalized version. Therefore I would suggest using the same sollution as in b)

d) Query 4 (same as 2)

This query will be doing the same as in query 2 but for the people without a subscription

```
MariaDB [Oblig5]> SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
-> FROM Car C JOIN Passing P USING(regno)
-> JOIN TaxClass T ON C.taxclass=T.idTaxClass
-> JOIN Fee F ON F.taxclass=T.idTaxClass
-> Left JOIN Subscription S USING(regno)
-> WHERE F.type='withsubscription'
-> AND S.regno IS NULL
-> GROUP BY C.owner HAVING totalfee > 4000;
```

```
+-----+-----+
| carowner | totalfee |
+-----+-----+
| Glen FjÃ_rtoft | 4010.00 |
| Lasse Nakken | 4316.00 |
| Ulva Hanssen | 4130.00 |
| Yulia Lie | 4260.00 |
+-----+-----+
```

4 rows in set (3.746 sec)

(difference being that we do a Left JOIN on Subscription and look for the regnr not in the subscription table.)

First lets have a look at the performance of this query as is.

```
MariaDB [Oblig5]> show profiles;
+-----+-----+-----+
| Query_ID | Duration | Query |
+-----+-----+-----+
| 1 | 3.72908210 | SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee |
| | | FROM Car C JOIN Passing P USING(regno) |
| | | JOIN TaxClass T ON C.taxclass=T.idTaxClass |
| | | JOIN Fee F ON F.taxclass=T.idTaxClass |
| | | Left JOIN Subscription S USING(regno) |
| | | WHERE F.type='withsubscription' |
| 2 | 3.70657548 | SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee |
| | | FROM Car C JOIN Passing P USING(regno) |
| | | JOIN TaxClass T ON C.taxclass=T.idTaxClass |
| | | JOIN Fee F ON F.taxclass=T.idTaxClass |
| | | Left JOIN Subscription S USING(regno) |
| | | WHERE F.type='withsubscription' |
| 3 | 3.62856547 | SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee |
| | | FROM Car C JOIN Passing P USING(regno) |
| | | JOIN TaxClass T ON C.taxclass=T.idTaxClass |
| | | JOIN Fee F ON F.taxclass=T.idTaxClass |
| | | Left JOIN Subscription S USING(regno) |
| | | WHERE F.type='withsubscription' |
| 4 | 3.70396787 | SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee |
| | | FROM Car C JOIN Passing P USING(regno) |
| | | JOIN TaxClass T ON C.taxclass=T.idTaxClass |
| | | JOIN Fee F ON F.taxclass=T.idTaxClass |
| | | Left JOIN Subscription S USING(regno) |
| | | WHERE F.type='withsubscription' |
| 5 | 3.71108078 | SELECT SQL_NO_CACHE C.owner AS carowner, Sum(F.costPerPassing) AS totalfee |
| | | FROM Car C JOIN Passing P USING(regno) |
| | | JOIN TaxClass T ON C.taxclass=T.idTaxClass |
| | | JOIN Fee F ON F.taxclass=T.idTaxClass |
| | | Left JOIN Subscription S USING(regno) |
| | | WHERE F.type='withsubscription' |
+-----+-----+-----+
5 rows in set (0.000 sec)
```

Now lets have a look at why these queries take almost 4s

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | T | index | PRIMARY | PRIMARY | 2 | NULL | 10 | Using index; Using |
| | | | | | | | | | |
| 1 | SIMPLE | F | ref | TeacherFK | TeacherFK | 2 | Oblig5.T.idTaxClass | 1 | Using where |
| | | | | | | | | | |
| 1 | SIMPLE | C | ref | PRIMARY,fk Car TaxClass1 | fk Car TaxClass1 | 2 | Oblig5.T.idTaxClass | 1 | |
| | | | | | | | | | |
| 1 | SIMPLE | S | eq ref | PRIMARY | PRIMARY | 7 | Oblig5.C.regno | 1 | Using where; Using index; |
| | | | | | | | | | |
| 1 | SIMPLE | P | ref | PRIMARY,reg | PRIMARY | 7 | Oblig5.C.regno | 13 | |
| | | | | | | | | | |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | T | index | PRIMARY | PRIMARY | 2 | NULL | 10 | Using index; Using |
| | | | | | | | | | |
| 1 | SIMPLE | F | ref | TeacherFK | TeacherFK | 2 | Oblig5.T.idTaxClass | 1 | Using where |
| | | | | | | | | | |
| 1 | SIMPLE | C | ref | PRIMARY,fk Car TaxClass1 | fk Car TaxClass1 | 2 | Oblig5.T.idTaxClass | 1 | |
| | | | | | | | | | |
| 1 | SIMPLE | S | eq ref | PRIMARY | PRIMARY | 7 | Oblig5.C.regno | 1 | Using where; Using index; |
| | | | | | | | | | |
| 1 | SIMPLE | P | ref | PRIMARY,reg | PRIMARY | 7 | Oblig5.C.regno | 13 | |
| | | | | | | | | | |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

Now this explain shows a similar story to the others. I don't see a Index help for this as it is using indexes and has available keys.

Will be testing the same normalization scheme as in **b)** as this is a very similar operation.

MariaDB [Oblig5]> show profiles;

Query_ID	Duration	Query
1	3.16445208	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='regular' GROUP BY owner HAVING totalfee > 4000
2	3.11762077	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='regular' GROUP BY owner HAVING totalfee > 4000
3	3.12460796	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='regular' GROUP BY owner HAVING totalfee > 4000
4	3.08037990	SELECT SQL_NO_CACHE owner AS carowner, Sum(costPerPassing) AS totalfee from Passing3 WHERE type='regular' GROUP BY owner HAVING totalfee > 4000

4 rows in set (0.000 sec)

The improvement is there however. 3.08 being the lowest is about .6s from the original query.

e) Query 5 (same as 3)

This query will be the same as Query 4 just using subqueries.

```
MariaDB [Oblig5]> SELECT SQL_NO_CACHE C.owner AS carowner,
-> Sum(F.costPerPassing) AS totalfee
-> FROM Car C JOIN Passing P USING(regno)
-> JOIN TaxClass T ON C.taxclass=T.idTaxClass
-> JOIN Fee F ON F.taxclass=T.idTaxClass
-> WHERE F.type='withsubscription'
-> AND C.regno NOT IN (SELECT regno FROM Subscription)
-> GROUP BY C.owner HAVING totalfee > 4000
-> ;
```

```
+-----+-----+
| carowner          | totalfee |
+-----+-----+
| Glen FjÃ¸rtoft    | 4010.00  |
| Lasse Nakken      | 4316.00  |
| Ulva Hanssen      | 4130.00  |
| Yulia Lie         | 4260.00  |
+-----+-----+
```

4 rows in set (3.746 sec)

(Difference being from «Query 3» being «not in» regno subscription)

```
MariaDB [Oblig5]> show profiles;
+-----+-----+
| Query_ID | Duration | Query |
+-----+-----+
|          |          |          |
+-----+-----+
|          |          |          |
+-----+-----+
| 1 | 3.78913586 | SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
FROM Car C JOIN Passing P USING(regno)
JOIN TaxClass T ON C.taxclass=T.idTaxClass
JOIN Fee F ON F.taxclass=T.idTaxClass
WHERE F.type='withsubscription'
AND C.regno NOT IN (SELECT regno FROM Subscription)
GROUP BY C.owner |
| 2 | 3.66125143 | SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
FROM Car C JOIN Passing P USING(regno)
JOIN TaxClass T ON C.taxclass=T.idTaxClass
JOIN Fee F ON F.taxclass=T.idTaxClass
WHERE F.type='withsubscription'
AND C.regno NOT IN (SELECT regno FROM Subscription)
GROUP BY C.owner |
| 3 | 3.67321359 | SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
FROM Car C JOIN Passing P USING(regno)
JOIN TaxClass T ON C.taxclass=T.idTaxClass
JOIN Fee F ON F.taxclass=T.idTaxClass
WHERE F.type='withsubscription'
AND C.regno NOT IN (SELECT regno FROM Subscription)
GROUP BY C.owner |
| 4 | 3.79695175 | SELECT SQL_NO_CACHE C.owner AS carowner,
Sum(F.costPerPassing) AS totalfee
FROM Car C JOIN Passing P USING(regno)
JOIN TaxClass T ON C.taxclass=T.idTaxClass
JOIN Fee F ON F.taxclass=T.idTaxClass
WHERE F.type='withsubscription'
AND C.regno NOT IN (SELECT regno FROM Subscription)
GROUP BY C.owner |
+-----+-----+
4 rows in set (0.000 sec)
```

This query did about as well as the previous one. And since it is the same operation the denormalization from the previous query can be used.

f) Query 6

Denne spørringen finner navnet til de som har gjort en passering på en spesifikk dag og tid.

```
MariaDB [Oblig5]> SELECT SQL_NO_CACHE C.owner FROM Car C WHERE C.regno
-> IN (SELECT P.regno FROM Passing P JOIN Tollstation T
-> ON P.tollstation = T.idTollstation
-> WHERE T.name LIKE 'Gravdal'
-> AND YEAR(P.timestamp)=2018
-> AND MONTH(P.timestamp)=2
-> AND DAYOFWEEK(P.timestamp)=1
-> AND HOUR(P.timestamp) = 3);
```

owner
Olav Håttta
Bjarnhild Reistad
Jarle Aarnes
Noor Evensen
Danny Straume
Hossein Kolberg
Hassan Haugland
Odin Persson
Enid Nicolaysen
Oddveig Roald
Xhavit Håiland
Vebjørn Bauge
Ramona Sævik
Oda Dammen
Nicholas Heiberg

--!!Shortening the output--

177 rows in set (3.396 sec)

```
MariaDB [Oblig5]> EXPLAIN SELECT SQL_NO_CACHE C.owner FROM Car C WHERE
C.regno
-> IN (SELECT P.regno FROM Passing P JOIN Tollstation T
-> ON P.tollstation = T.idTollstation
-> WHERE T.name LIKE 'Gravdal'
-> AND YEAR(P.timestamp)=2018
-> AND MONTH(P.timestamp)=2
-> AND DAYOFWEEK(P.timestamp)=1
-> AND HOUR(P.timestamp) = 3);
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	C	ALL	PRIMARY	NONE	NONE	NONE	203998	
1	PRIMARY	P	ref	PRIMARY, fk_Passing_1_reg	PRIMARY	7	Oblig5.C.regno	13	Using where
1	PRIMARY	T	eq_ref	PRIMARY	PRIMARY	2	Oblig5.P.tollstation	1	Using where; FirstMatch(C)

3 rows in set (0.001 sec)

Here we see there being a null under key. This is an indication that perhaps we need an index. But since it does have a key to use. And after hinting to use indexes it still didn't use it.

I don't believe much to be gained through denormalization but perhaps having Station name and owner in Passing would help.

Denormalization

Table:

```
CREATE TABLE IF NOT EXISTS Passing4 (
    regno CHAR(7) NOT NULL,
    timestamp TIMESTAMP NOT NULL,
    tollstation SMALLINT UNSIGNED NOT NULL,
    owner VARCHAR(85),
    name VARCHAR(45),
    PRIMARY KEY ( regno , timestamp ),
    CONSTRAINT fk_Passing_4
        FOREIGN KEY ( tollstation )
        REFERENCES Tollstation ( idTollstation )
        ON DELETE NO ACTION
        ON UPDATE NO ACTION,
    CONSTRAINT fk_Passing_Subscription4
        FOREIGN KEY ( regno )
        REFERENCES Car ( regno )
        ON DELETE NO ACTION
        ON UPDATE NO ACTION)
ENGINE = InnoDB;
```

Input data:

```
-- Passing4
INSERT INTO Passing4
SELECT DISTINCT regno,tid,idTollstation,OwnerName,tollname
FROM ImportTable;
```

New query:

```
MariaDB [Oblig5]> CREATE INDEX RegNr ON Passing (regno);

SELECT SQL_NO_CACHE owner FROM Passing4
WHERE name LIKE 'Gravdal'
AND YEAR(P.timestamp)=2018
AND MONTH(P.timestamp)=2
AND DAYOFWEEK(P.timestamp)=1
AND HOUR(P.timestamp) = 3);
```

Here we can also perhaps speed things up using a index on name.

Profiling

All queries are ran 10 times before profiling is started.

Pre Denormalization

MariaDB [Oblig5]> SHOW profiles;

Query_ID	Duration	Query
1	6.77211777	SELECT SQL_NO_CACHE C.owner FROM Car C WHERE C.regno IN (SELECT P.regno FROM Passing P JOIN Tollstation T ON P.tollstation = T.idTollstation WHERE T.name LIKE 'Gravdal' AND YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=2 AND DAYOFWEEK(P.timestamp)=1 AND HOUR(P.timestamp) = 3)
2	6.99386417	SELECT SQL_NO_CACHE C.owner FROM Car C WHERE C.regno IN (SELECT P.regno FROM Passing P JOIN Tollstation T ON P.tollstation = T.idTollstation WHERE T.name LIKE 'Gravdal' AND YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=2 AND DAYOFWEEK(P.timestamp)=1 AND HOUR(P.timestamp) = 3)
3	7.04556520	SELECT SQL_NO_CACHE C.owner FROM Car C WHERE C.regno IN (SELECT P.regno FROM Passing P JOIN Tollstation T ON P.tollstation = T.idTollstation WHERE T.name LIKE 'Gravdal' AND YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=2 AND DAYOFWEEK(P.timestamp)=1 AND HOUR(P.timestamp) = 3)
4	7.02457586	SELECT SQL_NO_CACHE C.owner FROM Car C WHERE C.regno IN (SELECT P.regno FROM Passing P JOIN Tollstation T ON P.tollstation = T.idTollstation WHERE T.name LIKE 'Gravdal' AND YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=2 AND DAYOFWEEK(P.timestamp)=1 AND HOUR(P.timestamp) = 3)
5	6.86503810	SELECT SQL_NO_CACHE C.owner FROM Car C WHERE C.regno IN (SELECT P.regno FROM Passing P JOIN Tollstation T ON P.tollstation = T.idTollstation WHERE T.name LIKE 'Gravdal' AND YEAR(P.timestamp)=2018 AND MONTH(P.timestamp)=2 AND DAYOFWEEK(P.timestamp)=1 AND HOUR(P.timestamp) = 3)

5 rows in set (0.000 sec)

Post denormalization

MariaDB [Oblig5]> show profiles;

Query_ID	Duration	Query
1	2.82387532	SELECT SQL_NO_CACHE owner FROM Passing4 WHERE name LIKE 'Gravdal' AND YEAR(timestamp)=2018 AND MONTH(timestamp)=2 AND DAYOFWEEK(timestamp)=1 AND HOUR(timestamp) = 3
2	2.82266015	SELECT SQL_NO_CACHE owner FROM Passing4 WHERE name LIKE 'Gravdal' AND YEAR(timestamp)=2018 AND MONTH(timestamp)=2 AND DAYOFWEEK(timestamp)=1 AND HOUR(timestamp) = 3
3	2.91836424	SELECT SQL_NO_CACHE owner FROM Passing4 WHERE name LIKE 'Gravdal' AND YEAR(timestamp)=2018 AND MONTH(timestamp)=2 AND DAYOFWEEK(timestamp)=1 AND HOUR(timestamp) = 3
4	2.88128561	SELECT SQL_NO_CACHE owner FROM Passing4 WHERE name LIKE 'Gravdal' AND YEAR(timestamp)=2018 AND MONTH(timestamp)=2 AND DAYOFWEEK(timestamp)=1 AND HOUR(timestamp) = 3
5	2.93795634	SELECT SQL_NO_CACHE owner FROM Passing4 WHERE name LIKE 'Gravdal' AND YEAR(timestamp)=2018 AND MONTH(timestamp)=2 AND DAYOFWEEK(timestamp)=1 AND HOUR(timestamp) = 3

5 rows in set (0.000 sec)

This did better than the original Query slicing the time spent in half.

Having a look at this new query with explain

```
MariaDB [Oblig5]> EXPLAIN SELECT SQL_NO_CACHE owner FROM Passing4
-> WHERE name LIKE 'Gravdal'
-> AND YEAR(timestamp)=2018
-> AND MONTH(timestamp)=2
-> AND DAYOFWEEK(timestamp)=1
-> AND HOUR(timestamp) = 3;
```

id	select_type	table	type	possible_keys	key	key_len	ref
1	SIMPLE	Passing4	ALL	NULL	NULL	NULL	NULL

5054935 Using where

1 row in set (0.000 sec)

We have no keys to use in this. Trying to create an index for «name»

```
MariaDB [Oblig5]> CREATE INDEX StationName ON Passing4(name);
```

The optimizer now uses the index created

```
+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows |
| Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | Passing4 | range | StationName | StationName | 48 | NULL | 352184 |
| Using index condition |
+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set (0.042 sec)
```

Checking if this helps for the performance of the query.

MariaDB [Oblig5]> show profiles;

```
+-----+-----+-----+-----+
| Query_ID | Duration | Query |
+-----+-----+-----+-----+
| 1 | 0.17989276 | SELECT SQL_NO_CACHE owner FROM Passing4 |
| | | WHERE name LIKE 'Gravdal' |
| | | AND YEAR(timestamp)=2018 |
| | | AND MONTH(timestamp)=2 |
| | | AND DAYOFWEEK(timestamp)=1 |
| | | AND HOUR(timestamp) = 3 |
| 2 | 0.17888898 | SELECT SQL_NO_CACHE owner FROM Passing4 |
| | | WHERE name LIKE 'Gravdal' |
| | | AND YEAR(timestamp)=2018 |
| | | AND MONTH(timestamp)=2 |
| | | AND DAYOFWEEK(timestamp)=1 |
| | | AND HOUR(timestamp) = 3 |
| 3 | 0.18819179 | SELECT SQL_NO_CACHE owner FROM Passing4 |
| | | WHERE name LIKE 'Gravdal' |
| | | AND YEAR(timestamp)=2018 |
| | | AND MONTH(timestamp)=2 |
| | | AND DAYOFWEEK(timestamp)=1 |
| | | AND HOUR(timestamp) = 3 |
| 4 | 0.17855339 | SELECT SQL_NO_CACHE owner FROM Passing4 |
| | | WHERE name LIKE 'Gravdal' |
| | | AND YEAR(timestamp)=2018 |
| | | AND MONTH(timestamp)=2 |
| | | AND DAYOFWEEK(timestamp)=1 |
| | | AND HOUR(timestamp) = 3 |
| 5 | 0.17951723 | SELECT SQL_NO_CACHE owner FROM Passing4 |
| | | WHERE name LIKE 'Gravdal' |
| | | AND YEAR(timestamp)=2018 |
| | | AND MONTH(timestamp)=2 |
| | | AND DAYOFWEEK(timestamp)=1 |
| | | AND HOUR(timestamp) = 3 |
+-----+-----+-----+-----+
5 rows in set (0.000 sec)
```

That is quite the improvement from the original query. From 6-7s to ~0.2s